

intro-to-data-science-rfm-analysis

March 28, 2020

```
[1]: # This Python 3 environment comes with many helpful analytics libraries
      ↳ installed
      # It is defined by the kaggle/python docker image: https://github.com/kaggle/
      ↳ docker-python
      # For example, here's several helpful packages to load in

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will list
↳ all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# Any results you write to the current directory are saved as output.
```

/kaggle/input/sample-sales-data/sales_data_sample.csv

0.1 ### RFM Analysis

RFM (Recency, Frequency, Monetary) analysis is a marketing technique used to determine quantitatively which customers are the best ones by examining *how recently a customer has purchased* (recency), *how often they purchase* (frequency), and *how much the customer spends* (monetary). Using RFM analysis, customers are assigned a ranking number of 1,2,3,4 (with 4 being highest) for each RFM parameter. The three scores together are referred to as an RFM “cell”. The data is sorted to determine which customers were the best customers in the past, with a cell ranking of 444 being ideal.

import the dataset first

```
[2]: df = pd.read_csv('/kaggle/input/sample-sales-data/sales_data_sample.csv',
      ↳ encoding='unicode_escape')
      display(df.head(5))
```

ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	\
-------------	-----------------	-----------	-----------------	-------	---

0	10107	30	95.70	2	2871.00
1	10121	34	81.35	5	2765.90
2	10134	41	94.74	2	3884.34
3	10145	45	83.26	6	3746.70
4	10159	49	100.00	14	5205.27

	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...	\
0	2/24/2003 0:00	Shipped	1	2	2003	...	
1	5/7/2003 0:00	Shipped	2	5	2003	...	
2	7/1/2003 0:00	Shipped	3	7	2003	...	
3	8/25/2003 0:00	Shipped	3	8	2003	...	
4	10/10/2003 0:00	Shipped	4	10	2003	...	

	ADDRESSLINE1	ADDRESSLINE2	CITY	STATE	\
0	897 Long Airport Avenue	NaN	NYC	NY	
1	59 rue de l'Abbaye	NaN	Reims	NaN	
2	27 rue du Colonel Pierre Avia	NaN	Paris	NaN	
3	78934 Hillside Dr.	NaN	Pasadena	CA	
4	7734 Strong St.	NaN	San Francisco	CA	

	POSTALCODE	COUNTRY	TERRITORY	CONTACTLASTNAME	CONTACTFIRSTNAME	DEALSIZE
0	10022	USA	NaN	Yu	Kwai	Small
1	51100	France	EMEA	Henriot	Paul	Small
2	75508	France	EMEA	Da Cunha	Daniel	Medium
3	90003	USA	NaN	Young	Julie	Medium
4	NaN	USA	NaN	Brown	Julie	Medium

[5 rows x 25 columns]

Let's filter the columns we need. we will only need four columns. **CUSTOMERNAME** to group customers, **ORDERDATE** to calculate recency, **ORDERNUMBER** to calculate frequency, and **SALES** to calculate monetary.

```
[3]: cols = ['CUSTOMERNAME', 'ORDERDATE', 'ORDERNUMBER', 'SALES']
df = df[cols]
print(df.head(5))
```

	CUSTOMERNAME	ORDERDATE	ORDERNUMBER	SALES
0	Land of Toys Inc.	2/24/2003 0:00	10107	2871.00
1	Reims Collectables	5/7/2003 0:00	10121	2765.90
2	Lyon Souvenirs	7/1/2003 0:00	10134	3884.34
3	Toys4GrownUps.com	8/25/2003 0:00	10145	3746.70
4	Corporate Gift Ideas Co.	10/10/2003 0:00	10159	5205.27

0.2 ### Recency:

To find recency we have to find the date of the last order of each customer made and subtract the value from the most recent order date in the dataset. The difference give us the recency value, for

recent customer the value will be smaller.

```
[4]: df['ORDERDATE'] = pd.to_datetime(df['ORDERDATE'])

#group the data by CUSTOMERNAME and only retrieve the ORDERDATE column
recent_order = df.groupby('CUSTOMERNAME')['ORDERDATE'].max()

most_recent = df['ORDERDATE'].max()

def subtract_date(date):
    days = (most_recent - date).days
    return days

recency = recent_order.apply(subtract_date)

print('Recency days : ', recency)
#print(recent_order.head(10))
#print(df.groupby(['CUSTOMERNAME', 'ORDERDATE']).count().head())
```

```
Recency days : CUSTOMERNAME
AV Stores, Co.      195
Alpha Cognac        64
Amica Models & Co.  264
Anna's Decorations, Ltd  83
Atelier graphique  187
...
Vida Sport, Ltd    274
Vitachrome Inc.    207
Volvo Model Replicas, Co  193
West Coast Collectables Co.  488
giftsbymail.co.uk  211
Name: ORDERDATE, Length: 92, dtype: int64
```

0.3 Frequency

Frequency is the measure how often a customer purchase a product. Now lets find out the number of times each customer has placed an order.

```
[5]: frequency = df.groupby(['CUSTOMERNAME', 'ORDERNUMBER']).size()
frequency = frequency.groupby('CUSTOMERNAME').size()
print(frequency.head())
```

```
CUSTOMERNAME
AV Stores, Co.      3
Alpha Cognac        3
Amica Models & Co.  2
Anna's Decorations, Ltd  4
Atelier graphique  3
dtype: int64
```

size() and count() function both same but count() ignores missing/NaN/NULL values.

0.4 ### Monetary

Monetary is how each customer spent.

```
[6]: #groupby the CUSTOMERNAME and only retrive the SALES and sum
monetary = df.groupby('CUSTOMERNAME')['SALES'].sum()
print(monetary.head())
```

```
CUSTOMERNAME
AV Stores, Co.          157807.81
Alpha Cognac           70488.44
Amica Models & Co.      94117.26
Anna's Decorations, Ltd 153996.13
Atelier graphique      24179.96
Name: SALES, dtype: float64
```

0.5 ### Putting All Together

Now we are going to put all these 3 parameter together.

```
[7]: rfm = pd.DataFrame()
rfm['recency'] = recency
rfm['frequency'] = frequency
rfm['monetary'] = monetary

print(rfm.head())
```

```
                recency  frequency  monetary
CUSTOMERNAME
AV Stores, Co.          195         3  157807.81
Alpha Cognac           64         3   70488.44
Amica Models & Co.      264         2   94117.26
Anna's Decorations, Ltd  83         4  153996.13
Atelier graphique      187         3   24179.96
```

Now we will convert these raw value into class, based on which quantile it fall into.

```
[8]: quantile_df = rfm.quantile([0.25,0.50,0.75])
display(quantile_df)
```

```
    recency  frequency  monetary
0.25    80.25         2.0   70129.4325
0.50   185.00         3.0   86522.6100
0.75   229.25         3.0  120575.8750
```

```
[9]: def quantile_classes(x, quantile_value, attribute):
    if attribute == 'recency':
```

```

        if x <= quantile_value.loc[0.25,attribute]: # receny is less than 0.25%
            return '4'
        elif x >= quantile_value.loc[0.25,attribute] and x <= quantile_value.
→loc[0.50,attribute]: # receny is larger than 25%
            return '3'
        elif x >= quantile_value.loc[0.50,attribute] and x <= quantile_value.
→loc[0.75,attribute]:
            return '2'
        else:
            return '1'
    else:
        #frequency and monetary
        if x <= quantile_value.loc[0.25,attribute]: # frequency/monetary is less
→than 0.25%
            return '1'
        elif x >= quantile_value.loc[0.25,attribute] and x <= quantile_value.
→loc[0.50,attribute]: # frequency/monetary is larger than 25%
            return '2'
        elif x >= quantile_value.loc[0.50,attribute] and x <= quantile_value.
→loc[0.75,attribute]:
            return '3'
        else:
            return '4'

#convert rfm table raw value to class
rfm['recency_class'] = rfm['recency'].apply(quantile_classes, args =
→(quantile_df, 'recency'))
rfm['frequency_class'] = rfm['frequency'].apply(quantile_classes, args =
→(quantile_df, 'frequency'))
rfm['monetary_class'] = rfm['monetary'].apply(quantile_classes, args =
→(quantile_df, 'monetary'))

display(rfm.head())

```

	recency	frequency	monetary	recency_class \
CUSTOMERNAME				
AV Stores, Co.	195	3	157807.81	2
Alpha Cognac	64	3	70488.44	4
Amica Models & Co.	264	2	94117.26	1
Anna's Decorations, Ltd	83	4	153996.13	3
Atelier graphique	187	3	24179.96	2

	frequency_class	monetary_class
CUSTOMERNAME		
AV Stores, Co.	2	4
Alpha Cognac	2	2

Amica Models & Co.	1	3
Anna's Decorations, Ltd	4	4
Atelier graphique	2	1

combine all of these individual class into a single column.

```
[10]: #join the string values
rfm['rfm_comb'] = rfm['recency_class'] + rfm['frequency_class'] +
    ↪rfm['monetary_class']

#convert to numeric value
rfm['rfm_comb'] = pd.to_numeric(rfm['rfm_comb'])

#sort values
rfm = rfm.sort_values(by=['rfm_comb'], ascending=False)

#display top 10 customer
display(rfm.head(10))
```

	recency	frequency	monetary	recency_class \
CUSTOMERNAME				
Salzburg Collectables	14	4	149798.63	4
Souveniers And Things Co.	2	4	151570.98	4
Mini Gifts Distributors Ltd.	2	17	654858.06	4
Danish Wholesale Imports	46	5	145041.60	4
Diecast Classics Inc.	1	4	122138.14	4
La Rochelle Gifts	0	4	180124.90	4
The Sharp Gifts Warehouse	39	4	160010.27	4
Reims Collectables	62	5	135042.94	4
Euro Shopping Channel	0	26	912294.11	4
Tokyo Collectables, Ltd	39	4	120562.74	4

	frequency_class	monetary_class	rfm_comb
CUSTOMERNAME			
Salzburg Collectables	4	4	444
Souveniers And Things Co.	4	4	444
Mini Gifts Distributors Ltd.	4	4	444
Danish Wholesale Imports	4	4	444
Diecast Classics Inc.	4	4	444
La Rochelle Gifts	4	4	444
The Sharp Gifts Warehouse	4	4	444
Reims Collectables	4	4	444
Euro Shopping Channel	4	4	444
Tokyo Collectables, Ltd	4	3	443